

PACEFUL

Technical Architecture & Methodology

Emotional readiness scoring infrastructure · v1.0 · May 2026

What Peaceful Is

Peaceful is API infrastructure for emotional readiness scoring. The platform accepts unstructured text inputs (chat messages, journal entries, customer interactions, claim notes, session transcripts) and returns structured signals across five operational dimensions, scored against vertical-specific signal packs.

Peaceful is positioned as an infrastructure layer, not a clinical product. Customers retain responsibility for clinical interpretation, regulatory compliance, and end-user workflow design.

The Five Dimensions

Each input is scored across five operational dimensions. These dimensions are designed for behavioral signal extraction from text, not for clinical diagnosis.

- **Emotional Stability** — consistency and regulation of affect over time within and across inputs.
- **Self-Reflection** — presence and depth of introspective language, attribution patterns, and metacognitive markers.
- **Coping Capacity** — evidence of adaptive strategies, resource utilization, and stress response patterns.
- **Behavioral Engagement** — indicators of action, follow-through, and goal-directed behavior.
- **Social Readiness** — relational language, social context awareness, and interpersonal orientation.

Each vertical signal pack maps these five dimensions to context-specific signals (e.g., the gambling pack maps to chase behavior, loss-of-control language, and recovery markers; the healthcare pack maps to depressive markers, treatment adherence indicators, and distress trajectory).

What Peaceful Does NOT Do

Explicit scope boundaries. These boundaries are intentional and structural, not aspirational.

- **No clinical diagnosis.** Peaceful does not diagnose, treat, or make clinical determinations. Output is signal data for downstream interpretation by qualified personnel.
 - **No replacement for human judgment.** Output is intended to inform decisions, not make them. Customer workflows must include human review for any consequential action.
 - **No FDA-regulated claims.** Peaceful is not currently registered as Software as a Medical Device. Customers who deploy in regulated contexts assume regulatory responsibility.
 - **No crisis intervention without explicit signal pack.** Crisis-detection signals (suicidal ideation, acute risk markers) are deliberately excluded from default packs and require explicit customer review and implementation before activation.
 - **No raw text retention.** Inputs are processed in-memory and discarded post-scoring. Peaceful does not store, log, or train on customer text.
 - **No PII extraction or storage.** Identifiable information passing through the API is not parsed, stored, or returned.
-

Technical Architecture

Request / Response Shape

All scoring requests are stateless POST operations. Authentication is via X-API-Key header. Default response time is sub-500ms p95 for inputs under 2,000 tokens.

Data Handling

- **In-memory processing.** Input text is held in memory only for the duration of the scoring request, then discarded.
- **No retention.** No raw text, derivative embeddings, or input artifacts are persisted to storage.
- **No model training.** Customer text is never used to train, fine-tune, or improve the scoring infrastructure.
- **Customer-controlled consent.** The customer is responsible for end-user consent. Peaceful does not collect end-user consent directly.
- **Aggregated, anonymized telemetry.** Score distributions, latency, and error rates are logged for service operation. No input content is logged.

Signal Pack Architecture

Verticals are configured via JSON signal pack definitions. A signal pack maps the five operational dimensions to vertical-specific named signals, defines threshold logic, and specifies output schema. Five productized packs are live: insurance, customer experience, gambling, dating, and healthcare.

Custom packs can be defined by enterprise customers for proprietary use cases. Peaceful retains pack format ownership; customers retain configuration ownership.

Scoring Engine

The scoring engine is proprietary infrastructure. Architecture details are available under NDA for evaluation purposes. The engine is designed for deterministic output stability across repeat queries on identical input, with versioned releases to ensure customer-side reproducibility.

Validation Roadmap

Peaceful is a v1 infrastructure product. The honest framing of where the platform stands today, and where it is going, is below.

Where We Are Today

- **Infrastructure stability.** 8 production endpoints, sub-500ms p95 latency, X-API-Key auth, in-memory processing, no retention. 10/10 tests passing on the Snapshot Assessment API.
- **Signal pack completeness.** Five productized verticals shipped with 64 total signals. Each pack is internally consistent and end-to-end tested.
- **Sandbox availability.** Public sandbox at peaceful-app.vercel.app/sandbox, no signup required, supports paste-test of all five verticals.
- **Output stability.** Deterministic scoring across repeat queries on identical input, validated through internal regression testing.

Where We Are Going (Next 12 Months)

- **Clinical advisory board.** Recruitment in progress. Healthcare pack signals (particularly any future crisis-detection signals) will require advisory sign-off prior to release.
- **Design partner co-development.** Five anchor customers will co-develop next-generation signal packs. Their domain expertise informs vertical evolution; their deployments inform real-world validation.
- **Construct alignment study.** Mapping of Peaceful's five operational dimensions to validated psychological constructs (Big Five, established clinical instruments) under research collaboration. Findings will be published as a methodology white paper.
- **Third-party validation.** External validation studies planned for healthcare pack as customer base scales.
- **Compliance posture.** SOC 2 Type II in 2026. HIPAA Business Associate Agreement (BAA) capability available for healthcare design partners.

Honest Version Status

Peaceful v1 is infrastructure. The platform has been engineered for stability, latency, and structural correctness. It has not yet undergone third-party clinical validation. Customers deploying in healthcare or other regulated

contexts should treat Peaceful as a signal infrastructure provider and apply their own clinical, regulatory, and safety review prior to production use.

This positioning is intentional. We would rather ship honest infrastructure that survives technical due diligence than overclaim clinical depth we have not yet validated.

Working with Peaceful

Design Partner Program

Five design partner slots open for 2026. \$7,200 / month, 12-month commitment. Partners receive: priority signal pack co-development, direct founder access, custom signal definition, and methodology input.

Evaluation Path

- **Sandbox.** peaceful-app.vercel.app/sandbox — no signup, paste-test any of five verticals.
- **Technical evaluation.** API key issued for time-bound evaluation; full endpoint access; sample integration provided.
- **Pilot deployment.** Production-grade deployment with monitoring, support SLA, and dedicated channel.

Contact

LJ — Founder, Peaceful
partners@peaceful.com
peaceful-app.vercel.app

Peaceful — Emotional readiness scoring as infrastructure.